



Questions and Answers

Topics that reviewers found unclear

asked by Thistle (pseudonym)

recorded on 2015-04-10

What control will a user have over wiping of files? Wiping of files is not practical? You mean manual wiping?

There are two different functions, regular deletion and permanent deletion. One is reversible and other is highly restricted. Please refer to answer to next question.

Practicality is meant here as both performance and assurance. Wiping in general is a synonym to permanent removal but in this context it means a specific implementation, through overwriting of file content. Historically, this is how it was done. There is at least one major deficiency with this approach, namely performance. Overriding a 4 GB file for example, even at hard-disks maximum 120 MB/s, still takes 34.1 seconds. And to be really sure, you should do that 35 times (Gutmann method).

Will there be an option for non-permanent wiping in case a user accidentally wipes a file? What safety measures will be there to prevent accidental complete wiping of content? Or is that a risk users have to take?

There are two distinct functions. Regular deletion (here called deletion) is an action that can be taken by non-privileged user. Rogue processes would be able to take this action without user interaction. Therefore this kind of action should be inspect-able and undo-able. Browsing history would reveal any deletion and allow for recovery. Permanent deletion (here called wiping) is a different action that requires root privileges. Rogue processes would not be able to take this action. This kind of action could also be inspect-able, as revision history would reveal that some wiping has taken place but file content would not be accessible anymore.

History of changes to files means that all versions of the file remain stored, or at least a history that can be viewed, isn't that impractical in some cases of confidential files? An option to not save a file history?

History of changes is a list of snapshots, states saved at some points in time, ordered chronologically. History can be browsed and every snapshot can be inspected for details but also every snapshot can be restored back. Confidentiality is not compromised because browsing past versions of files requires same level of access as browsing current version. Anyone having access to past history would necessarily need to have access to current file in the first place. When user deletes a file, it will be recoverable as long as it is browsable in history. When user permanently deletes a file (as opposed to regularly deleting), it will not be recoverable even if history mentions that such file existed in the past. Versioning can also be disabled.

Deleted files are permanently gone, how would that be done? In comparison to saving file history...

Technically there are two ways of ensuring data is permanently gone, either through overriding data itself or overriding encryption keys that were used to encrypt said data.

Second approach is much better since keys occupy only few bytes and overriding few bytes takes almost no time. Versioning is an independent feature.

Legal questions, may use of the program encounter difficulties with existing laws in some countries?

Depending on your country of residence, you may be:

- forbidden from using cryptographic products in general
- required to get license or send notification before importing (Russia)
- required to disclose cryptographic keys to authorities in advance (Russia)
- asked to decrypt your laptop contents (crossing US border)
- asked if you have any other undisclosed encrypted partitions (crossing US border)
- subpoenaed to produce keys or decrypted data itself
- jailed for long time or until you produce encryption keys (US and UK)

Lesson to learn here is that established law can put you in a position where technical solutions do not give you an easy and legal way out. Consider crossing US border with child pornography on your laptop (there was a real case). Officer asks you to mount all partitions and then asks if you have any hidden partitions on your laptop. Note that lying to a federal agent is a criminal offense in US. If you mount them all, jail, if you lie, possible jail, if you deny, some other consequences. Plausible deniability feature may let you lie out of it but it will not make it any more legal.

asked by Steven Balderrama

recorded on 2015-08-14

On mission document, page 3, I found top 3 things what lacks today in a filesystem: versioning, secure deletion and disk utilisation.

There are also other features but some are so specialized that clearly there will not be many people that will have a use for them. This project aims to deliver a versatile filesystem that will be useful to everybody.

But selective compression is great. Some files should not be compressed for performance.

Some files are almost impossible to compress, audio video formats are a good example. However, performance does not necessarily have to be negatively impacted. Methods like Gzip are meant to achieve high utilisation with little regard for performance but there are also other, less efficient, methods that are many times faster than a hard disk. Snappy developed by Google can compress ~250 MB/s and decompress ~500 MB/s per core at it's slowest while a hard disk can sustain ~120 MB/s transfer. Therefore processor load should be fractional and not impact user experience.

We sorta talked about performance. Since there is compression and encryption, how will that not hinder performance?

Compare throughput of modern hard disks to throughput of modern processors. Hard disks can sustain ~120 MB/s. Compression like Snappy can handle ~250 MB/s per core easily. Encryption using Salsa20 can handle ~400 MB/s per core. Lightweight hashing using SipHash can handle ~530 MB/s per core. You can see that all of the above can be achieved at a fraction of CPU load, and you can always disable some features to lower

the load. Clearly the hard disk is the bottleneck here and not the CPU.

And yes, versioning. How many times I have seen this, especially programmers who do not utilize software versioning. I have seen it even at my work now. Version control, is that what you mean?

There are two features that may be used to revert changes. Snapshots is a feature that saves state at explicitly chosen points in time, allowing to revert only to specific states that have been recognized in advance. This is the model of how version control software works. Commits are then the equivalent of creating snapshots. So this was the first feature. Continuous versioning is similar but has an important difference. Every change is automatically creating a temporary lightweight snapshot that gets stacked on top of previous changes. At the same time, changes from too far into the past are becoming non-revertible. So in contrast, snapshots are manually created while versioning is automatic, and snapshots are kept forever while versions are only available for some period of time.

Lastly, you said about having it's own type of recovery. Explain in detail how it will recover or roll back.

Early filesystems like Fast File System and it's descendant Ext2 depended on running a program called *fsck* that scanned all inodes after an interruption. Later filesystems like Ext3 were using journaling to reapply changed blocks. Later filesystems like Btrfs and ZetabyteFS started using copy on write and checksums. The state of the art solution would perhaps comprise of copy on write, checksums, and intents.

From user's perspective one thing should be noticeable: recovery should happen lazily. Mounting an interrupted filesystem should be achievable in sub-second time. Files that were not closed properly require maintenance and will be scrutinized either at opening or during routine scrubbing, whichever comes first. Maintenance should be carried out on individual files and be postponed preferably until after booting is complete. Further recovery can run in background. Unaffected files become accessible immediately.

This can be achieved by using copy on write, checksums, and intents. For example, when a write operation occurs the data is stored into a newly allocated extent and both its address and checksum are stored in an intent, which itself is stored at a known location. At mount only intents need to be processed. Intents integrity are checked through their checksums and if they are good, extents they point to become reserved as if they were properly allocated. At this point filesystem is mounted and ready to go. When a file gets opened any remaining unverified intents referring to it must be processed. Intent holds a checksum to verify the data, if data is good then intent is reapplied to inode, if not then data is deallocated and operation gets reverted.